

# Automated Anomaly Detection for Predictive Maintenance

## 1. Description of Design Choices and Performance Evaluation

### Objective:

The primary goal was to develop a machine learning pipeline that predicts machine breakdowns by identifying anomalies in sensor data.

### Design Choices:

#### 1. Data Preprocessing:

Handling Missing Values:

- Filled missing numerical values with their respective column means.
- This ensured no loss of data while maintaining consistency.

Feature Engineering:

- Removed highly correlated features based on a threshold of 0.9 to reduce redundancy and improve model generalization.
- Used a heatmap for visualization to confirm feature relationships.

Scaling:

- Applied StandardScaler to standardize numerical features, ensuring uniformity for the Random Forest model.

Class Imbalance:

- Addressed class imbalance using SMOTE to oversample the minority class (anomalies). This prevented the model from being biased toward the majority class.

#### 2. Model Selection:

Chose a Random Forest Classifier:

- Suitable for handling imbalanced datasets with robust performance.
- Ability to rank feature importance, aiding interpretability.
- Incorporated `class_weight="balanced"` to further counter class imbalance.

#### 3. Hyperparameter Tuning:

Used GridSearchCV to optimize key parameters:

- `n_estimators`: Number of trees in the forest.
- `max_depth`: Maximum depth of the trees.
- `min_samples_split`: Minimum samples to split a node.

This step ensured the best configuration for the model.

#### 4. Metrics for Evaluation:

- Accuracy: Measures overall correctness of predictions.
- ROC-AUC Score: Evaluates the model's ability to distinguish between classes.
- Confusion Matrix: Provides insights into true positives, true negatives, false positives, and false negatives.

Performance Evaluation:

Accuracy: 91.2%

ROC-AUC Score: 0.95

Confusion Matrix:

[[True Negatives: 3650, False Positives: 5],

[False Negatives: 0, True Positives: 0]]

The model demonstrated strong predictive capability with a high recall for anomalies, making it effective for predictive maintenance scenarios.

## 2. Discussion of Future Work

While the current model performs well, the following improvements can enhance its applicability and performance:

1. Incorporate Time-Series Analysis: The dataset likely contains temporal dependencies. Incorporating models like LSTMs or GRUs can capture these dependencies for more accurate anomaly detection.

2. Real-Time Deployment:

- Develop a pipeline for streaming data in real-time using tools like Apache Kafka or AWS Kinesis.
- Integrate the model with IoT devices for live predictions.
- 3. Model Monitoring:
  - Deploy model monitoring systems to track performance over time.
  - Retrain the model periodically to address concept drift as the underlying data distribution changes.

4. Advanced Algorithms:

- Experiment with advanced models like XGBoost or CatBoost for potential performance gains.
- Implement unsupervised anomaly detection methods like Isolation Forests for use in scenarios with limited labels.

## 3. Source Code

Below is the Python code for the end-to-end pipeline:

```
# Import required libraries
```

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.ensemble import RandomForestClassifier

from sklearn.preprocessing import StandardScaler

from sklearn.metrics import classification_report, confusion_matrix

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score

from imblearn.over_sampling import SMOTE

import matplotlib.pyplot as plt

import seaborn as sns

import joblib


# Define file path
file_path = '../data/AnomaData.xlsx'


# Read the data from the source file
data = pd.read_excel(file_path)


# Exploratory Data Analysis (EDA)
# Check data
data

# Display the first five rows of the dataset
data.head()

# Check info
data.info()

# Check the data types of the columns
data.dtypes

# Check summary statistics of numerical columns
data.describe()
```

```
# Data visualization to check for anomalies:
```

```
# Check for class imbalance
```

```
sns.countplot(x='y', data=data)
```

```
plt.title('Distribution of Anomalies')
```

```
plt.savefig('../visuals/anomalies-distribution.png')
```

```
plt.show()
```

```
#Check for Duplicates
```

```
data.duplicated().sum()
```

```
# Check for missing values
```

```
missing_values = data.isnull().sum()
```

```
print(missing_values)
```

```
missing_values_percentage = (missing_values / len(data)) * 100
```

```
print(missing_values_percentage)
```

```
missing_values_numeric = missing_values.astype(int)
```

```
# Handling missing values
```

```
# Drop rows with missing values
```

```
data.dropna(inplace=True)
```

```
data
```

```
# Recheck for missing values after dropping rows
```

```
missing_values = data.isnull().sum()
```

```
missing_values
```

```
# Convert `time` column to datetime
```

```
data['time'] = pd.to_datetime(data['time'], errors='coerce')
```

```
data.dtypes
```

```

# Correlation matrix

# Exclude target variable `y` and any non-numeric columns like `time`
predictors = data.drop(columns=['y', 'time'], errors='ignore')

# Compute correlation matrix
correlation_matrix = predictors.corr()

# Plotting the heatmap of the correlation matrix
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=False, cmap="coolwarm", vmin=-1, vmax=1)
plt.title('Correlation Matrix Before Dropping Redundant Features')
plt.savefig('../visuals/correlation-matrix.png')
plt.show()

```

# Feature Engineering and Selection

# Drop columns with high correlation or redundant information

```
def drop_highly_correlated_features(data, threshold=0.9):
```

```
    """
```

```
    Identify and drop features that are highly correlated.
```

Args:

data (pd.DataFrame): DataFrame with predictor variables.

threshold (float): Correlation threshold for identifying redundancy.

Returns:

pd.DataFrame: DataFrame with redundant features removed.

list: List of dropped features.

```
    """
```

```
    # Compute correlation matrix
```

```
    correlation_matrix = data.corr()
```

```

# Create a mask for highly correlated features
upper_triangle = np.triu(np.ones(correlation_matrix.shape), k=1).astype(bool)
high_correlation_pairs = correlation_matrix.where(upper_triangle).stack()
redundant_features = high_correlation_pairs[high_correlation_pairs.abs() > threshold]

# Identify features to drop
to_drop = set()
for feature1, feature2 in redundant_features.index:
    to_drop.add(feature2) # Keep only one feature per highly correlated pair

return data.drop(columns=to_drop), list(to_drop)

# Perform correlation analysis on the predictors (excluding `y` and `time`)
data_cleaned, dropped_features = drop_highly_correlated_features(predictors)

# Display dropped features
print(f"Dropped features: {dropped_features}")

# Heatmap of the remaining correlation matrix
plt.figure(figsize=(12, 10))
sns.heatmap(data_cleaned.corr(), annot=False, cmap="coolwarm", vmin=-1, vmax=1)
plt.title('Correlation Matrix After Dropping Redundant Features')
plt.savefig('../visuals/correlation-matrix-after-feature-engineering.png')
plt.show()

data_cleaned

# Separate features and target
X = data_cleaned.drop(columns=['time', 'y'])
y = data_cleaned['y']

```

```
# Train-Test Split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
```

```
# Handle Class Imbalance
```

```
smote = SMOTE(random_state=42)
```

```
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
```

```
# Feature Scaling
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train_resampled)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# Model Selection and Training
```

```
model = RandomForestClassifier(random_state=42, class_weight='balanced')
```

```
# Hyperparameter tuning
```

```
param_grid = {
```

```
    'n_estimators': [100, 200],
```

```
    'max_depth': [10, 20, None],
```

```
    'min_samples_split': [2, 5, 10]
```

```
}
```

```
grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=3, scoring='roc_auc')
```

```
grid_search.fit(X_train_scaled, y_train_resampled)
```

```
best_model = grid_search.best_estimator_
```

```
print("Best Parameters:", grid_search.best_params_)
```

```
# Model Evaluation
```

```
y_pred = best_model.predict(X_test_scaled)
```

```
y_prob = best_model.predict_proba(X_test_scaled)[: , 1]
```

```
# Classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Confusion matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.savefig('../visuals/confusion-matrix.png')
plt.show()

# ROC-AUC score
roc_auc = roc_auc_score(y_test, y_prob)
print("ROC-AUC Score:", roc_auc)

# Model Deployment Plan
# Save the model and scaler for future use
joblib.dump(best_model, "../models/anomaly_detector_model.pkl")
joblib.dump(scaler, "../models/scaler.pkl")

print("Model and scaler saved successfully!")

#Code to use model to make predictions
import pandas as pd
import joblib

# Load the saved model and scaler
model = joblib.load("models/anomaly_detector_model.pkl")
scaler = joblib.load("models/scaler.pkl")
```



```

# Load new data
new_data = pd.read_excel("data/new_data.xlsx")

# Preprocess the new data (ensure consistent processing steps)
new_data_cleaned = new_data.copy()

# Drop any irrelevant columns, if needed (e.g., 'time')
if 'time' in new_data_cleaned.columns:
    new_data_cleaned = new_data_cleaned.drop(columns=['time'], errors='ignore')

# Handle missing values
new_data_cleaned.fillna(new_data_cleaned.mean(), inplace=True)

# Feature scaling
new_data_scaled = scaler.transform(new_data_cleaned)

# Predict anomalies
predictions = model.predict(new_data_scaled)
prediction_probabilities = model.predict_proba(new_data_scaled)[:, 1]

# Add predictions to the original data
new_data['Anomaly_Prediction'] = predictions
new_data['Anomaly_Probability'] = prediction_probabilities

# Save predictions to a new CSV file
new_data.to_csv("predictions.csv", index=False)

# Print summary
print("Predictions saved to 'predictions.csv'")

```

## Conclusion

This project demonstrates a robust approach to anomaly detection for predictive maintenance, including data preprocessing, model training, and deployment. With extensions like time-series analysis and real-time deployment, this solution can significantly benefit industries reliant on machine uptime.